

# How to use Celery+RabbitMQ

Nick Thompson

November 22, 2015

## Getting started:

```
$ git clone https://github.com/NAThompson/learn_celery.git  
$ cd learn_celery  
$ pip3 install -r requirements.txt
```

## What will we attempt to accomplish?

- ▶ Link Celery+RabbitMQ into Django ecosystem

## What will we attempt to accomplish?

- ▶ Link Celery+RabbitMQ into Django ecosystem
- ▶ Use tasks to keep page loads from timing out

## What will we attempt to accomplish?

- ▶ Link Celery+RabbitMQ into Django ecosystem
- ▶ Use tasks to keep page loads from timing out
- ▶ Associate a running job with a progress bar

## What will we attempt to accomplish?

- ▶ Link Celery+RabbitMQ into Django ecosystem
- ▶ Use tasks to keep page loads from timing out
- ▶ Associate a running job with a progress bar
- ▶ Link CPU cycles to billing

# What will we attempt to accomplish?

- ▶ Link Celery+RabbitMQ into Django ecosystem
- ▶ Use tasks to keep page loads from timing out
- ▶ Associate a running job with a progress bar
- ▶ Link CPU cycles to billing
- ▶ Autoscale nodes based on queue length

# Why Celery + RabbitMQ

- ▶ Job scheduling is boring, and it's a pain.



# Why Celery + RabbitMQ

- ▶ Job scheduling is boring, and it's a pain.
- ▶ So you certainly don't want to learn *two* job schedulers

# Why Celery + RabbitMQ

- ▶ Job scheduling is boring, and it's a pain.
- ▶ So you certainly don't want to learn *two* job schedulers
- ▶ You never want to be in a spot where you have to swap out your job scheduler

# Why Celery + RabbitMQ

- ▶ Job scheduling is boring, and it's a pain.
- ▶ So you certainly don't want to learn *two* job schedulers
- ▶ You never want to be in a spot where you have to swap out your job scheduler
- ▶ Therefore, use the most well-tested and popular stack

# What is RabbitMQ?

- ▶ RabbitMQ is an implementation of the *advanced messaged queuing protocol* (AMQP), a standardized way to pass data between applications.
- ▶ The specification of AMQP is [300 pages long](#), but it's basically an email server protocol for binary applications

## Warning!

- ▶ Celery as an API is a bit baffling

## Warning!

- ▶ Celery as an API is a bit baffling
- ▶ Many people contend that using RabbitMQ directly leads to less astonishment

## Warning!

- ▶ Celery as an API is a bit baffling
- ▶ Many people contend that using RabbitMQ directly leads to less astonishment
- ▶ The author has never tried RabbitMQ directly, but has been astonished many times by Celery. So keep this in mind as an option.

## How do I set up RabbitMQ?

```
$ sudo apt-get install -y rabbitmq-server  
$ sudo rabbitmqctl status  
$ sudo lsof -i :5672 # 5672 is the well-known port of AMQP
```

If your server still isn't up, or you need to restart it:

```
$ sudo rabbitmq-server -detached
```



## Testing RabbitMQ

We want to get the minimal working RabbitMQ working on localhost:

```
$ rabbitmqctl list_queues
```

```
Listing queues ...
```

```
$ python3 basic_rabbitmq/send_message.py
```

```
We have sent the message 'Hello World' to RabbitMQ
```

```
$ rabbitmqctl list_queues
```

```
Listing queues ...
```

```
hello 1
```

```
$ python3 basic_rabbitmq/send_message.py
```

```
We have sent the message 'Hello World' to RabbitMQ
```

```
$ rabbitmqctl list_queues
```

```
Listing queues ...
```

```
hello 2
```

## Receive the messages:

```
learn_celery$ python3 basic_rabbitmq/receive_message.py
Waiting for messages. To exit press CTRL+C
[x] Received (b'Hello World',)
[x] Received$ (b'Hello World',)
learn_celery rabbitmqctl list_queues
Listing queues ...
hello 0
```

## Deleting a queue

As you learn RabbitMQ, you start generating meaningless queues that you might want to get rid of. Here's how to delete them:

```
learn_celery$ rabbitmqctl list_queues
Listing queues ...
hello 0
(learn_celery) ~/learn_celery$ rabbitmqctl stop_app
Stopping node 'rabbit@nthompson-Precision-M6700' ...
(learn_celery) ~/learn_celery$ rabbitmqctl reset
Resetting node 'rabbit@nthompson-Precision-M6700' ...
(learn_celery) ~/learn_celery$ rabbitmqctl start_app
Starting node 'rabbit@nthompson-Precision-M6700' ...
(learn_celery) ~/learn_celery$ rabbitmqctl list_queues
Listing queues ...
```

## On to Celery

- ▶ You could connect to your RabbitMQ directly, without Celery, but there's a lot of boilerplate.
- ▶ Celery makes a nice glue between your Django and Rabbitmq.

## Celery Basic Example

```
$ cd basic_celery
$ celery --app=task work --loglevel=info &
$ python3 -q
>>> from tasks import add
>>> r = add.delay(12, 12)
[2015-11-05 14:16:07,806: INFO/MainProcess] Received task: tasks.add[4f
[2015-11-05 14:16:07,816: INFO/MainProcess] Task tasks.add[4f2c482f-5e6
>>> r.status
'SUCCESS'
>>> r.result
24
```

Easy peasy.

## Celery Basic Example

Celery will start some queues in RabbitMQ after startup:

```
# rabbitmqctl list_queues
```

```
Listing queues ...
```

```
celery0
```

```
celery@nthompson-Precision-M6700.celery.pidbox0
```

```
celeryev.39a8ba96-9ed0-4b31-83a3-0728cb6273c50
```

## Multiple workers

Celery has support for multiple workers:

```
celery --app=tasks worker --loglevel=INFO --concurrency=nproc
----- celery@nthompson-Precision-M6700 v3.1.19 (Cipater)
---- **** -----
--- * *** * -- Linux-3.19.0-32-generic-x86_64-with-Ubuntu-15.04-vivid
-- * - **** ---
- ** ----- [config]
- ** ----- .> app:          tasks:0x7feabb181128
- ** ----- .> transport:    amqp://guest:**@localhost:5672//
- ** ----- .> results:      rpc://
- *** --- * --- .> concurrency: 8 (prefork)
-- ***** ---
--- ***** ----- [queues]
----- .> celery          exchange=celery(direct) key=celery
```

The default concurrency is nproc.

# Embedding Celery into Django

It's not a huge jump from this to embedding in Django:

- Specify the serialization types you'll accept in settings.py:

```
CELERY_ACCEPT_CONTENT = ['json']  
CELERY_TASK_SERIALIZER = 'json'  
CELERY_RESULT_SERIALIZER = 'json'
```



# Embedding Celery into Django

It's not a huge jump from this to embedding in Django:

- ▶ Specify the serialization types you'll accept in settings.py:

```
CELERY_ACCEPT_CONTENT = ['json']  
CELERY_TASK_SERIALIZER = 'json'  
CELERY_RESULT_SERIALIZER = 'json'
```

- ▶ Specify the broker url in settings.py

```
BROKER_URL = 'amqp://guest:guest@localhost//'
```

# Embedding Celery into Django

- ▶ Start a app and add some tasks to it:

```
$ ./manage.py startapp jobs
```

```
$ touch jobs/tasks.py # Define a task in here. Must be named tasks.
```

# Embedding Celery into Django

- ▶ Start a app and add some tasks to it:

```
$ ./manage.py startapp jobs
```

```
$ touch jobs/tasks.py # Define a task in here. Must be named tasks.
```

- ▶ Add your app to INSTALLED\_APPS in settings.py, and

```
echo "from .celery import app as celery_app # noqa" >> __init__.py
```

# Embedding Celery into Django

- ▶ Start a app and add some tasks to it:

```
$ ./manage.py startapp jobs
```

```
$ touch jobs/tasks.py # Define a task in here. Must be named tasks.
```

- ▶ Add your app to INSTALLED\_APPS in settings.py, and

```
echo "from .celery import app as celery_app # noqa" >> __init__.py
```

- ▶ Now start up celery:

```
$ celery --app=celery_in_django worker --loglevel=INFO &
```

You should see your task registered in the logs.

## Minimal Django+Celery Example:

Assuming you have RabbitMQ up, then:

```
learn_celery$ git checkout basic_async_task
```

```
learn_celery$ cd django_example
```

```
django_example$ celery --app=celery_in_django worker --loglevel=info &
```

```
django_example$ ./manage.py runserver
```

Now navigate to 127.0.0.1:8000, and start hitting buttons.

# Not-as-minimal Django+Celery Example

Problems with the minimal example:

- ▶ The user gets no feedback when he's pressed a button; as such might press it 40 times

# Not-as-minimal Django+Celery Example

Problems with the minimal example:

- ▶ The user gets no feedback when he's pressed a button; as such might press it 40 times
- ▶ The user doesn't know when the task is done

# Not-as-minimal Django+Celery Example

Problems with the minimal example:

- ▶ The user gets no feedback when he's pressed a button; as such might press it 40 times
- ▶ The user doesn't know when the task is done
- ▶ The user doesn't know how much to estimate how long the job will run



## Take 1: Store tasks in the session

This is a low-effort, effective for some tasks method:

- ▶ Store task id in `request.session`; then use the task id as an argument to `celery.result.AsyncResult` constructor.

## Take 1: Store tasks in the session

This is a low-effort, effective for some tasks method:

- ▶ Store task id in request.session; then use the task id as an argument to celery.result.AsyncResult constructor.
- ▶ Then we make buttons available based on task status:

```
$ git checkout task_in_session  
$ cd django_example  
$ ./manage.py runserver
```

## Take 1: Store tasks in the session

Note that by default, the task state moves from 'PENDING' directly to 'FINISHED'.

This can be remedied by putting the following in settings.py:

```
CELERY_TRACK_STARTED = True
```

## Take 1: Store tasks in the session

- ▶ This is not a great design, as the session should expire at browser close, and we might want to run jobs for longer than a session.

## Take 1: Store tasks in the session

- ▶ This is not a great design, as the session should expire at browser close, and we might want to run jobs for longer than a session.
- ▶ Hence we need to store task id's in our database.

## Take 1: Store tasks in the session

- ▶ This is not a great design, as the session should expire at browser close, and we might want to run jobs for longer than a session.
- ▶ Hence we need to store task id's in our database.
- ▶ Note: This will require user logins, the code is going to get more complex!

## Take 2: Store task in the Database

Once we have user logins, we can associate each user with a list of tasks.

```
$ git checkout tasks_in_db
$ basic_celery/kill_celery.sh
$ cd django_example
$ ./start_celery.sh
$ ./manage.py runserver
```

This will show us what time we started the task, as well as show when the task is started and finished (but no timestamps yet).

# Progress Reporting

An accepted method of doing progress reporting in Celery is the use of [custom states](#). Let's look at the code:

```
$ git checkout progress_bar
$ basic_celery/kill_celery.sh
$ cd django_example
$ ./start_celery.sh
$ ./manage.py runserver
```



## Progress Bars

If the progress bar is moving too slowly, we might want to terminate the running job. Easy: Link it to a post request, user the celery “revoke” function:

```
$ git checkout revoke
$ basic_celery/kill_celery.sh
$ cd django_example
$ ./start_celery.sh
$ ./manage.py runserver
```

## Job metadata

Note that the job metadata is being printed to stdout by Celery:

```
[2015-11-22 19:13:35,024: INFO/MainProcess]  
Task jobs.tasks.fft_random[123e5a8d-1cda-457c-9082-810d4f163572]  
succeeded in 13.250278769002762s: 0.20759175911474792
```